

FYS-2021: Assignment 1

TOR-IVAR HASSFJORD

UiT - Norges Arktiske Universitet

08.09.2024

Abstract

This report explores how logistic regression can be applied to classify Spotify songs into either Pop or Classical using stochastic gradient descent (SGD). We focus on two key features: *liveness* and *loudness*. Throughout this report, I attempt to apply what I learned in class, such as using SGD, feature scaling, and evaluating models using metrics like accuracy, precision, recall, and F1 score.

Contents

1	Method/Theory	2
1.1	Logistic Regression	2
1.2	Stochastic Gradient Descent (SGD)	2
1.3	Feature Scaling	2
1.4	Shuffling the Data	2
1.5	Overfitting	3
1.6	Confusion Matrix and Evaluation Metrics	3
1.7	Connection to Course Content	3
2	Implementation	3
2.1	Task 1: Data Preprocessing	3
2.1.1	1a) Loading the Data	3
2.1.2	1b) Shuffling the Data	4
2.1.3	1c) Train-Test Split	4
2.1.4	1d) Feature Scaling	4
2.2	Task 2: Model Training	4
2.2.1	2a) Implementing Logistic Regression with SGD	4
2.2.2	2b) Training the Model	5
2.2.3	Task 2c: Visualizing the Decision Boundary	5
2.3	Task 3: Model Evaluation	5
2.3.1	3a) Confusion Matrix	5
3	Results and Discussion	6
3.1	Task 1 Results: Data Preprocessing	6
3.2	Task 2 Results: Model Training	6
3.3	Task 3 Results: Model Evaluation	8
3.3.1	Evaluation Metrics	8
3.4	Training vs Test Accuracy: Overfitting Analysis	8
3.5	Impact of Feature Scaling	9
3.6	Overfitting and Generalization	9
3.7	Model Complexity vs Overfitting Tradeoff	9

Introduction

In this assignment, I was asked to classify songs from the Spotify dataset into Pop or Classical genres using logistic regression. I decided to use stochastic gradient descent (SGD) as the optimization algorithm, as this was required by the assignment. I used two features: *liveness* and *loudness*. The model's performance was evaluated using confusion matrices, and several metrics like accuracy, precision, recall, and F1 score. In the sections that follow, I'll walk through the methods, implementation, and results.

1 Method/Theory

In this section, I want to explain some of the key concepts that I learned in my course and how they apply to the work I did in this report. I'll go through how logistic regression works, why we use stochastic gradient descent (SGD), the importance of feature scaling, and the metrics used to evaluate the model.

1.1 Logistic Regression

Logistic regression is used when we want to classify data into two categories, which is what I am trying to do here with Pop and Classical songs. The idea is to fit a line that best separates the two classes. This line comes from the equation:

$$\text{sigmoid}(z) = \frac{1}{1 + e^{-z}}$$

This sigmoid function outputs a value between 0 and 1, which we interpret as the probability of the data belonging to a certain class (in this case, Pop). Logistic regression is useful here because it's simple and interpretable, but I learned it's only good for problems where the classes are linearly separable, which might not be the case here.

1.2 Stochastic Gradient Descent (SGD)

The assignment asked us to use SGD to train our logistic regression model. SGD is an optimization method that updates the model's weights after each training sample, rather than waiting for the entire dataset (like regular gradient descent). I learned from the course slides that this makes it faster, especially with large datasets, but it's also noisier because it's based on individual samples.

The learning rate is an important part of SGD. It controls how big of a step we take in the direction of the gradient. In my findings, I observed that a learning rate of 0.0001 combined with 15 epochs resulted in the most optimal performance for this logistic regression model. This combination provided a good balance between training time and accuracy.

When I experimented with lower learning rates, such as 0.00001, the model still reached a similar accuracy, but it required significantly more epochs to achieve the same result. This extended training time without any noticeable improvement in performance made using lower learning rates inefficient and not worth the additional time.

1.3 Feature Scaling

I used a custom version of the standard scaler to normalize my features before feeding them into the model. This was important because features like loudness and liveness are on very different scales, and logistic regression assumes that features are scaled similarly. Without scaling, one feature might dominate the others. The 'CustomStandardScaler' I used ensures that the features are centered around 0 and have unit variance, which is a common preprocessing step for models like logistic regression.

1.4 Shuffling the Data

Before splitting the data into training and test sets, I shuffled the entire dataset. This was an important step to make sure the model doesn't learn any biases from the order of the data. I also shuffled the data before training to ensure the model wasn't influenced by any patterns that could exist in the order of the samples.

1.5 Overfitting

One thing we talked about in class is the problem of overfitting. Overfitting happens when the model performs really well on the training data but poorly on the test data, meaning it learned the "noise" in the training data instead of the actual patterns. I saw this happening a bit in my results because the training accuracy was a bit higher than the test accuracy (92.10% vs. 91.85%).

Overfitting is common when there's not enough data, or when the model is too complex, but in my case, logistic regression is a pretty simple model, so maybe it's because my features (loudness and liveness) don't completely capture the difference between Pop and Classical. In the slides, they suggested using cross-validation or adding more features to help with overfitting, so that might be something to try next time.

1.6 Confusion Matrix and Evaluation Metrics

The confusion matrix helps us see where the model is making mistakes. In my case, it looks like the model is better at identifying Pop songs (high true positives) than Classical ones, but it sometimes misclassifies Classical songs as Pop. I used a few different metrics to understand this better:

$$\begin{aligned}\text{Accuracy} &= \frac{TP + TN}{TP + TN + FP + FN} \\ \text{Precision} &= \frac{TP}{TP + FP} \\ \text{Recall} &= \frac{TP}{TP + FN} \\ \text{F1 Score} &= 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}\end{aligned}$$

Accuracy tells us how often the model was right overall, but precision and recall give us more insight into how it handles the two classes. Precision is about how many of the predicted Pop songs were actually Pop, while recall tells us how many of the real Pop songs were correctly identified. The F1 score balances precision and recall.

In my case, I had a recall of 98.47%, meaning the model found most of the Pop songs, but precision was lower at 87.16%, meaning it predicted some Classical songs as Pop. The F1 score of 92.47% shows the model strikes a decent balance between these two.

1.7 Connection to Course Content

Throughout this report, I tried to apply what we learned in class. The use of SGD, understanding overfitting, and using evaluation metrics like precision and recall all connect back to the core machine learning concepts we covered. I learned that machine learning is not just about building models, but also about understanding how and why they work, and how to evaluate them properly. This project helped me see how things like learning rates and feature scaling, which we talked about in the optimization and gradient descent lectures, make a big difference in practice.

2 Implementation

2.1 Task 1: Data Preprocessing

2.1.1 1a) Loading the Data

```
1 import pandas as pd
2 # Load the dataset
3 data = pd.read_csv('../data/SpotifyFeatures.csv')
4 pop_classical_df = data[data['genre'].isin(['Pop', 'Classical'])].copy()
5 pop_classical_df.loc[:, 'label'] = pop_classical_df['genre'].apply(lambda x: 1 if x == 'Pop' else 0)
```

The dataset contains 232725 songs with 18 features. After filtering for Pop and Classical, I ended up with 9386 Pop songs and 9256 Classical songs.

2.1.2 1b) Shuffling the Data

Before splitting the data into training and test sets, I shuffled the entire dataset to prevent any biases from the order of the songs.

```
1 import numpy as np
2 # Shuffle before splitting
3 np.random.seed(42)
4 shuffled_indices = np.random.permutation(len(pop_classical_df))
5 data_shuffled = pop_classical_df.iloc[shuffled_indices]
```

2.1.3 1c) Train-Test Split

I split the data into 80% for training and 20% for testing.

```
1 train_ratio = 0.8
2 # Split the data
3 train_size = int(train_ratio * len(data_shuffled))
4 X_train = features_shuffled[:train_size]
5 y_train = labels_shuffled[:train_size]
6 X_test = features_shuffled[train_size:]
7 y_test = labels_shuffled[train_size:]
```

2.1.4 1d) Feature Scaling

To ensure that both features were on the same scale, I used a custom standard scaler.

```
1 class CustomStandardScaler:
2     def fit(self, X):
3         self.mean_ = np.mean(X, axis=0)
4         self.std_ = np.std(X, axis=0)
5     def transform(self, X):
6         return (X - self.mean_) / self.std_
7     def fit_transform(self, X):
8         self.fit(X)
9         return self.transform(X)
10
11 scaler = CustomStandardScaler()
12 X_train_scaled = scaler.fit_transform(X_train)
13 X_test_scaled = scaler.transform(X_test)
```

2.2 Task 2: Model Training

2.2.1 2a) Implementing Logistic Regression with SGD

To train the logistic regression model, I used stochastic gradient descent.

```
1 def sigmoid(z):
2     return 1 / (1 + np.exp(-z))
3
4 def logistic_sgd(X, y, learning_rate, epochs):
5     weights = np.zeros(X.shape[1])
6     errors = []
7
8     for epoch in range(epochs):
9         for i in range(X.shape[0]):
10             z = X[i] @ weights
11             h = sigmoid(z)
12             gradient = (h - y[i]) * X[i]
13             weights -= learning_rate * gradient
14
15         y_pred = predict(X, weights)
16         error = np.mean(y_pred != y)
17         errors.append(error)
18
19     return weights, errors
```

2.2.2 2b) Training the Model

I trained the model using a learning rate of 0.0001 and ran it for 20 epochs.

```
1 weights, errors = logistic_sgd(X_train_scaled, y_train, learning_rate=0.0001, epochs=20)
```

2.2.3 Task 2c: Visualizing the Decision Boundary

For the bonus task, I extracted the learned parameters from the logistic regression model to draw the decision boundary. The plot below shows how the logistic regression model separates the two classes, Pop and Classical, based on the features *liveness* and *loudness*.

The green line represents the decision boundary learned by the logistic regression model, while the red and blue dots represent Pop and Classical songs, respectively.




Figure 1: Decision Boundary from Logistic Regression for Pop and Classical Songs

As seen in the plot (Figure 1), the logistic regression classifier tries to draw a linear boundary between the two classes based on the input features. While it performs well for distinguishing a large portion of the songs, it may struggle in areas where there is overlap between the classes. This could explain the slightly lower precision, as some Classical songs may be misclassified as Pop.

This plot helps visualize why the model performs the way it does. The separation between Pop and Classical songs isn't perfectly clear-cut, which is likely why the classifier isn't 100% accurate. Adding more features or using a more complex model might help better capture the separation.

2.3 Task 3: Model Evaluation

2.3.1 3a) Confusion Matrix

After training the model, I computed the confusion matrix to evaluate its performance.

```
1 def confusion_matrix(y_true, y_pred):
2     TP = np.sum((y_true == 1) & (y_pred == 1))
3     TN = np.sum((y_true == 0) & (y_pred == 0))
4     FP = np.sum((y_true == 0) & (y_pred == 1))
5     FN = np.sum((y_true == 1) & (y_pred == 0))
6     return np.array([[TP, FP], [FN, TN]])
7
8 conf_matrix = confusion_matrix(y_test, y_test_pred)
9 print(conf_matrix)
```

3 Results and Discussion

3.1 Task 1 Results: Data Preprocessing

The dataset contains 9386 Pop songs and 9256 Classical songs. After splitting and scaling, the features were successfully normalized, ensuring that both *liveness* and *loudness* had zero mean and unit variance. This step was crucial because logistic regression is sensitive to feature scales, and without scaling, the larger feature (loudness) would have dominated the model training process.

3.2 Task 2 Results: Model Training

The logistic regression model was trained using Stochastic Gradient Descent (SGD) with a learning rate of 0.0001 and 15 epochs. I experimented with different learning rates, and I found that this combination produced the best results. At this learning rate, the model converged efficiently without overshooting the optimal solution.

Training Error vs Epochs (SGD with Matrix Multiplication)




Figure 2: Training Error vs Epochs for Learning Rate 0.0001 and 15 Epochs (SGD with Matrix Multiplication)

As shown in Figure 2, the training error decreases significantly during the early epochs and stabilizes after about 10 epochs. This indicates that the model reaches a good solution quickly, and further training doesn't result in significant improvements.

When I experimented with higher learning rates, such as 0.001, I noticed instability in the early epochs, as shown in Figure 3. The model seemed a bit "jumpy" with a higher learning rate, and I think this happened because it was taking too big steps when trying to adjust the weights. This caused it to miss the best values, so the training error didn't go down as smoothly.




Figure 3: Training Error vs Epochs for Higher Learning Rate (SGD with Matrix Multiplication)

I also experimented with a lower learning rate (0.00001) but I had to increase the number of epochs to 50. As shown in Figure 4, the training error decreased more smoothly over the longer training period. However, it took significantly more epochs to reach a similar level of performance as the 15-epoch case. The model converged, but it required more time to train.




Figure 4: Training Error vs Epochs for Learning Rate 0.0001 with 50 Epochs (SGD with Matrix Multiplication)

Conclusion:

Based on my experiments, a learning rate of 0.0001 with 15 epochs seemed to work best. It allowed the model to train quickly without causing any major issues with accuracy or making the training unstable. When I tried using lower learning rates, the model was more stable, but I had to run it for a lot more epochs to get similar results, which just seemed inefficient. On the other hand, using higher learning rates made the model jump around too much during training, causing the error to fluctuate a lot.

3.3 Task 3 Results: Model Evaluation

The confusion matrix for the model on the test data is shown in Figure 5:




Figure 5: Confusion Matrix of the Logistic Regression Model

$$\begin{bmatrix} 1866 & 274 \\ 30 & 1559 \end{bmatrix}$$

From this matrix, we can observe that the model performs well in identifying Pop songs (high true positives, low false negatives), but it has a slightly higher rate of misclassifying Classical songs as Pop (274 false positives).

3.3.1 Evaluation Metrics

The following evaluation metrics were computed based on the confusion matrix:

- **Training Accuracy:** 92.10%
- **Test Accuracy:** 91.85%
- **Precision:** 87.16%
- **Recall (Sensitivity):** 98.47%
- **F1 Score:** 92.47%

These metrics show that the model is very effective at identifying Pop songs (high recall), but its precision is slightly lower because of the false positives (misclassified Classical songs). The F1 Score, which balances precision and recall, is still high, indicating that the model performs well overall.

3.4 Training vs Test Accuracy: Overfitting Analysis

In class, we talked about overfitting, which is when a model does really well on the training data but struggles to perform on new, unseen data. Overfitting is not good because it means the model is basically memorizing the training data rather than learning patterns that can generalize to other datasets.

In my case, the training accuracy came out to be 92.10%, and the test accuracy was very close at 91.85%. The small

difference between these two numbers suggests that the model might not be overfitting and is probably generalizing well. I think this is partly because logistic regression is a relatively simple model, and from what I've learned, simpler models tend to be less prone to overfitting compared to more complex ones. Another big factor could be the feature scaling I did. Before scaling, the model's accuracy was around 70% (+/- 10%), but after scaling the features, the accuracy jumped to over 90%. This shows how important scaling was for my model because it ensured that both *liveness* and *loudness* were treated equally when the model was learning.

Overall, the small gap between the training and test accuracy is a good sign. It indicates that the model is learning patterns that work for both the training data and new data, rather than just memorizing the training examples. I can't say for sure without further testing, but this gives me confidence that overfitting isn't a major issue here.

3.5 Impact of Feature Scaling

As I mentioned earlier, I applied feature scaling to both *liveness* and *loudness* to ensure that they were on the same scale. This was really important because, without scaling, the logistic regression model could have given more weight to the feature with the larger range (probably *loudness*), which would've led to poor performance. By scaling, I made sure that both features contributed equally to the learning process. This helped the model focus on finding a good decision boundary between Pop and Classical songs.

The scaling process had a massive impact on the model's performance. Before scaling, the model was stuck at around 70% accuracy, but after scaling, it jumped to over 90%. So, from what I've seen, scaling features is a crucial step, especially for models like logistic regression that are sensitive to the scale of input features.

3.6 Overfitting and Generalization

Based on the close match between training accuracy (92.10%) and test accuracy (91.85%), I think it's safe to say that the model generalizes well and isn't overfitting to the training data. However, there is a slight drop in performance from training to test accuracy, which could be explained by the limited number of features I used—just *liveness* and *loudness*. It's possible that adding more features, like *danceability* or *energy*, could help improve the model's ability to distinguish between Pop and Classical songs even further. That's something I'd like to explore if I had more time for experimentation.

3.7 Model Complexity vs Overfitting Tradeoff

One thing we talked about in class is the tradeoff between model complexity and overfitting. If a model is too simple, it might underfit the data, meaning it won't be able to capture the important patterns in the dataset, which is called having high bias. On the other hand, if the model is too complex, it might overfit the data, meaning it learns not just the patterns but also the noise and outliers, leading to poor generalization on new data (this is called having high variance).

In my case, I used logistic regression, which is a relatively simple model. Based on what I've learned, simpler models like this are less prone to overfitting, but they also run the risk of underfitting if the data is too complex or if important features are missing. For example, I only used *liveness* and *loudness* as features in my model. These features alone might not fully capture the differences between Pop and Classical songs, which could cause the model to underfit if there are more complex patterns in the data that it's not able to detect.

That said, the close match between training and test accuracy in my results suggests that the model is not overfitting, and it's generalizing well to new data. This makes sense because logistic regression tends to have high bias, meaning it's more likely to underfit than overfit, especially when we're working with a limited number of features. In the future, it would be interesting to see if adding more features or using a more complex model (like a decision tree or a neural network) would improve performance, but that also comes with the risk of overfitting.